

Continual Reinforcement Learning for Robotic Tasks

Initial Presentaiton

Philemon Schöpf

Supervisors: Antonio Rodriguez Sanchez, Sayantan Auddy

2022-03-15

Reinforecment Learning

- Learn by interacting with an environment
- Find a policy that takes the optimal action for each state (input)

Reinforcement Learning

- Learn by interacting with an environment
- Find a policy that takes the optimal action for each state (input)
- Ideally suitable to robotics
 - No (labeled) training data required
 - Physical simulation of environment easier

Continual Learning

- Learning of multiple tasks in succession
- Adapt existing model with new experiences
 - Model has to retain old information
 - Old training data is **not accessible** while training new tasks
- Critically important for human intelligence
- Still a major issue in machine learning

Catastrophic Forgetting

- Old skills are forgotten as network is trained on new data
 - No incentive to keep old skills
- Traditional approach: *replay* of old training examples to incentivize preservation
 - Keep old training data: storage inefficient, privacy issues
 - Not scalable to “lifelong learning”

Previous Work

- Regularization: penalize changes in prediction


$$\theta_s^*, \theta_o^*, \theta_n^* = \operatorname{argmin}_{\hat{\theta}_s, \hat{\theta}_o, \hat{\theta}_n} (\lambda \mathcal{L}_{old}(\mathbf{Y}_o, \hat{\mathbf{Y}}_o) + \mathcal{L}_{old}(Y_n, \hat{Y}_n) + \mathcal{R})$$

- Elastic Weight Consolidation (EWC): penalty according to Fisher information matrix

$$\mathcal{L} = \mathcal{L}_0 + \sum_i \frac{\lambda}{2} F_i (\theta_i - \theta_{A,i})^2$$

- Dynamic architectures: change/expand network while training
- Hypernetworks: NN that output another NN

- Simulated environment for door opening
- Different types of door handles
 - Round turn knob
 - Lever knob
 - Pull knob
- Train agent able to open all 3 kinds of doors
 - One type of door learned at a time

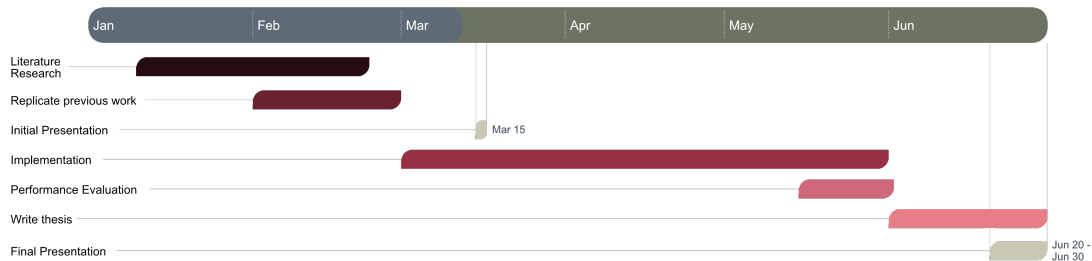


DoorGym demo

Goals

- Use the hypernetwork approach for reinforcement learning
- First: task-incremental CL
 - Type of door is known at inference time
- Expanding: domain- and class-incremental CL
 - Domain-incremental CL: type is unknown
 - Class-incremental CL: type is unknown *and* we want to infer type
- Inference e.g. via object recognition from renders

Planned Timeline



References



R. S. Sutton, A. G. Barto, Reinforcement Learning: An Introduction, 2nd Edition, The MIT Press, 2018.

URL <http://incompleteideas.net/book/the-book-2nd.html>



G. I. Parisi, R. Kemker, J. L. Part, C. Kanan, S. Wermter, Continual lifelong learning with neural networks: A review, CoRR abs/1802.07569 (2018).

URL <http://arxiv.org/abs/1802.07569>



Y. Huang, K. Xie, H. Bharadhwaj, F. Shkurti, Continual model-based reinforcement learning with hypernetworks, CoRR abs/2009.11997 (2020).

URL <https://arxiv.org/abs/2009.11997>



G. M. van de Ven, A. S. Tolias, Three scenarios for continual learning, CoRR abs/1904.07734 (2019).

URL <http://arxiv.org/abs/1904.07734>